



Data Curation Pipeline

Filtering Results Report

May 16, 2025

Analysis of 1 GitHub Repositories



Generated with SWE-RL Inspired Data Curation Pipeline

Executive Summary

This report presents the results of a data curation pipeline designed to extract and filter high-quality software engineering data from GitHub repositories.

Across 1 repositories, a total of 10 PRs were processed through the filtering pipeline, resulting in 5 high-quality PRs that passed all filters. This represents an overall pass rate of 50.0%, with a data reduction ratio of 50.0%.

Filtering Breakdown:

- Bot Filter: 1 PRs (10.0% of total)
- Size/Complexity Filter: 4 PRs (40.0% of total)
- Content Relevance Filter: 0 PRs (0.0% of total)

The average quality score for passing PRs was 0.82 on a scale of 0-1.

Key Findings:

1. Bot-generated PRs constitute a significant portion of repository activity
2. Size and complexity filters effectively remove both trivial and unwieldy changes
3. Content relevance filtering ensures focus on meaningful software engineering content
4. The pipeline successfully identifies related files that provide context for changes

The following pages provide detailed analyses for each repository and highlight exemplary PRs that represent high-quality software engineering data.

Filter Rates Across Repositories

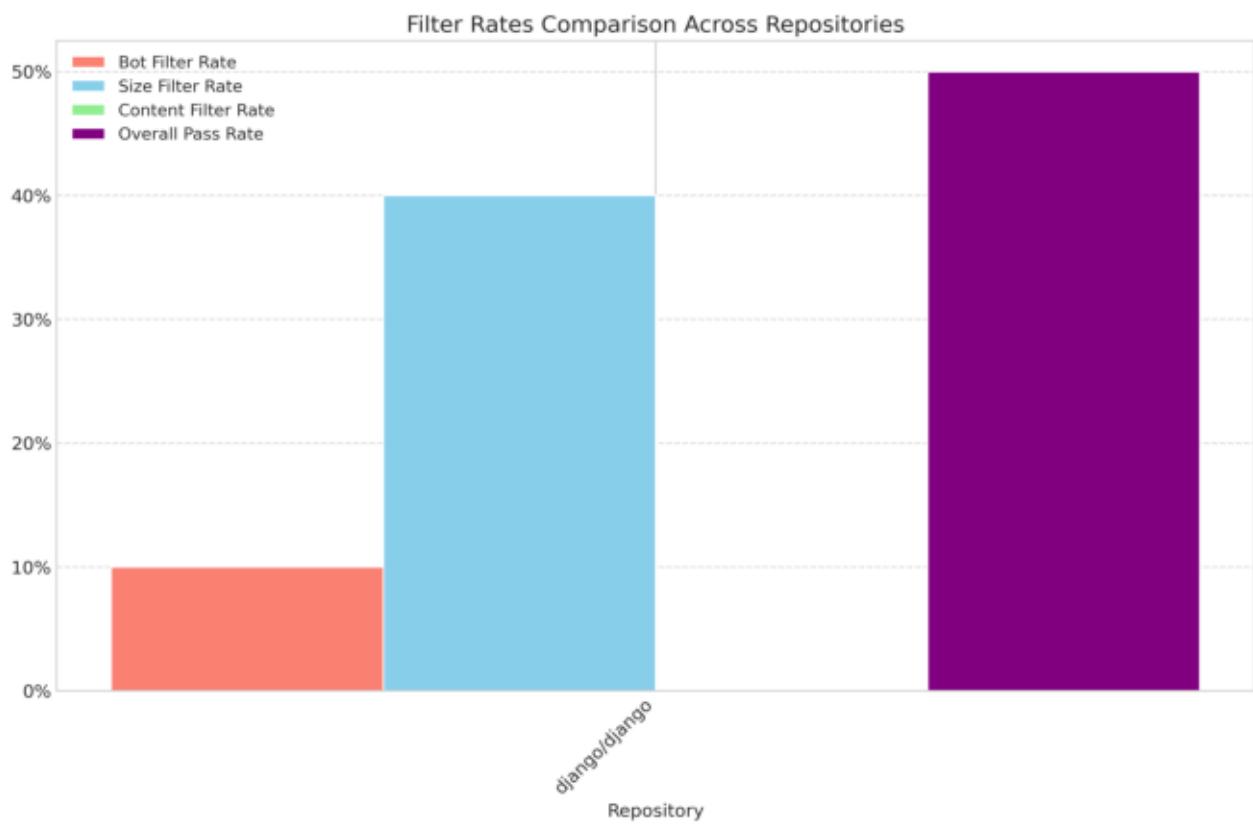


Figure: Filter Rates Across Repositories

Data Reduction Comparison




Figure: Data Reduction Comparison

Cross-Repository Comparison

Quality Metrics Comparison



Figure: Quality Metrics Comparison Across Repositories

Cross-Repository Metrics Comparison

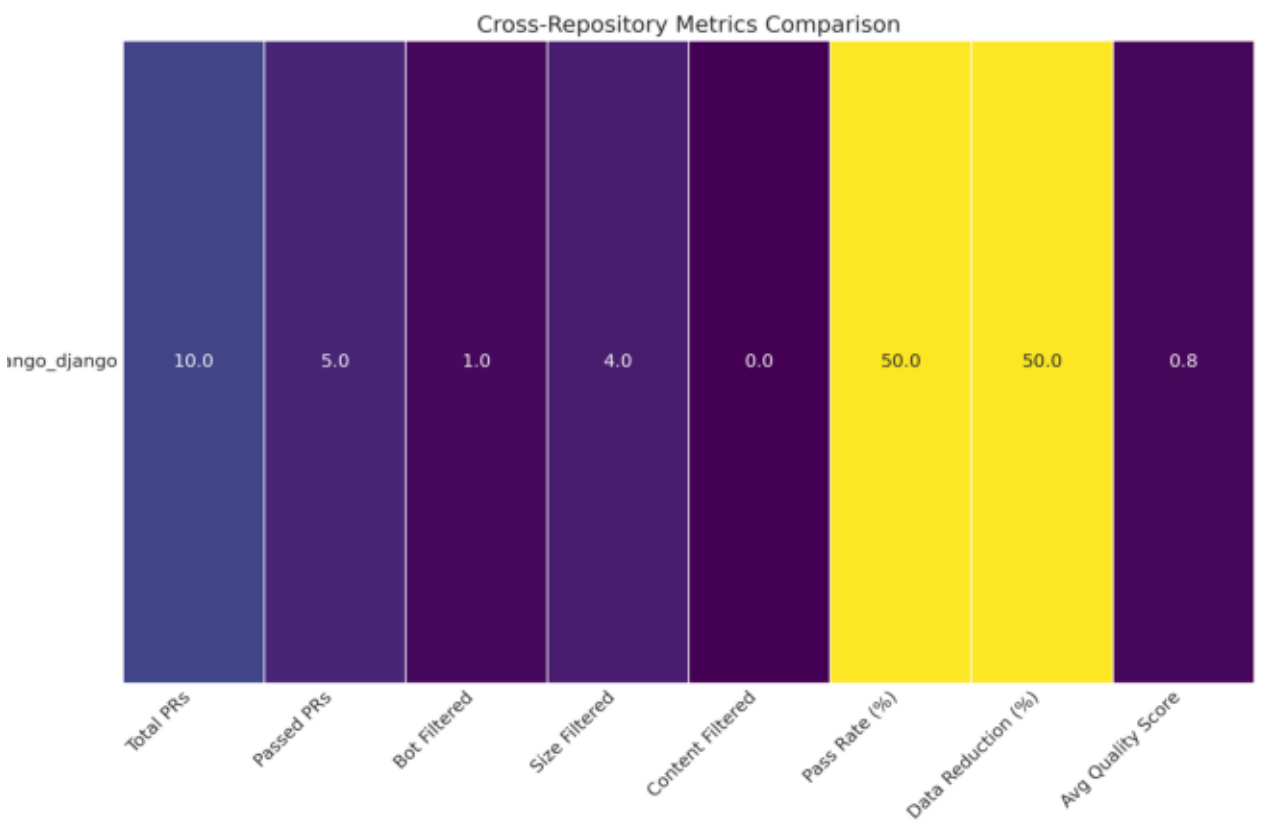


Figure: Cross-Repository Metrics Comparison

Repository: django/django

Total PRs: 10
Passed PRs: 5 (50.0%)
Data Reduction: 50.0%
Average Quality Score: 0.82

- Filtering Breakdown:
- Bot Filter: 1 PRs (10.0%)
 - Size Filter: 4 PRs (40.0%)
 - Content Filter: 0 PRs (0.0%)




Figure: Filter Funnel Analysis

Quality Distribution: django/django

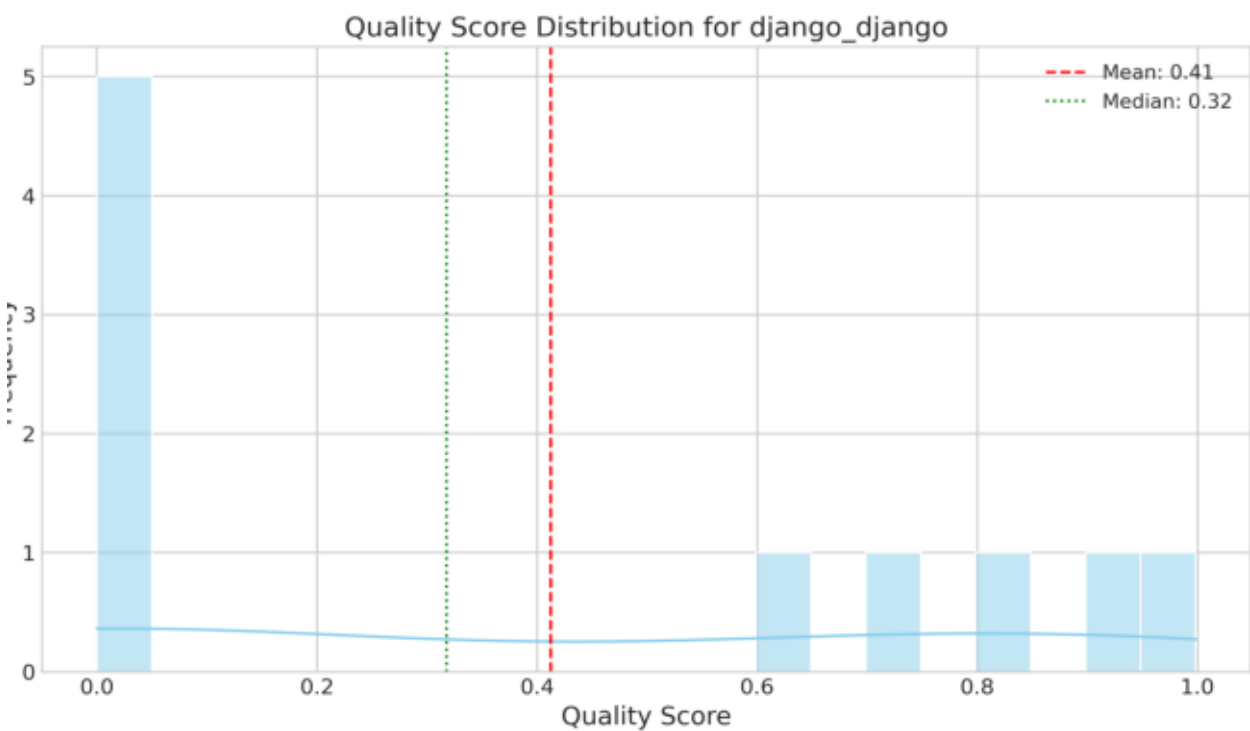



Figure: Quality Distribution: django/django

Filter Analysis: django/django

Bot Filter Reasons for django_django



Size Filter Rejection Reasons for django_django




Figure: Size Filter Stats

Quality Profiles of Exemplary PRs


The following pages showcase high-quality PRs that passed all filtering stages. These PRs represent exemplary software engineering data with meaningful problem-solving content, appropriate size, and high relevance scores.

Each scorecard provides detailed metrics on the PR's quality dimensions, including file composition, code changes, and identified relevant files that provide context for understanding the changes.

PR: [django/django/pr19469](#)






PR: [django/django/pr19419](#)



Detailed PR quality scorecards are presented on the following pages

PR Quality Scorecard: django/django/pr19469




PR #19469 Quality Scorecard - django_django



- Relevant Files
- Relevant Files (5 total)**
- django/db/models/sql/where.py
 - django/db/models/sql/datastructures.py
 - django/utils/regex_helper.py
 - django/utils/functional.py
 - django/db/models/fields/__init__.py

PR Quality Scorecard: django/django/pr19419

PR #19419 Quality Scorecard - django_django





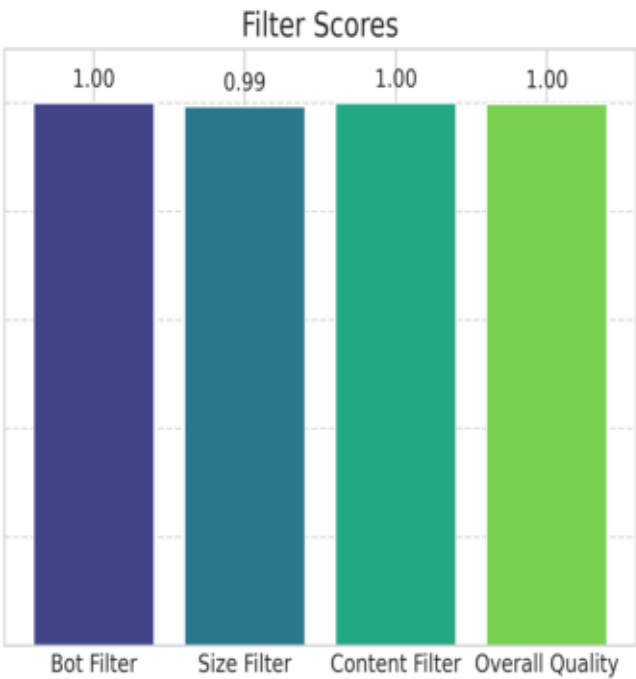
Relevant Files

Relevant Files (5 total)

- django/utils/formats.py
- django/utils/translation/__init__.py
- django/utils/safestring.py
- django/utils/timezone.py
- django/template/__init__.py

PR Quality Scorecard: django/django/pr19445

PR #19445 Quality Scorecard - django_django




- Relevant Files
- Relevant Files (5 total)**
- django/db/models/fields/composite.py
 - django/db/models/expressions.py
 - django/db/models/utils.py
 - tests/model_regress/test_state.py
 - django/contrib/contenttypes/fields.py

PR Quality Scorecard: django/django/pr19445

PR #19445 Quality Scorecard - django_django

The following pages showcase high-quality PRs that passed all filtering stages. These PRs represent exemplary software engineering data with meaningful problem-solving content, appropriate size, and high relevance scores.




Each scorecard provides detailed metrics on the PR's quality dimensions, including file composition, code changes, and identified relevant files that provide context for understanding the changes.



Detailed PR quality scorecards are presented on the following pages

PR Quality Scorecard: django/django/pr19469




PR #19469 Quality Scorecard - django_django



- Relevant Files
- Relevant Files (5 total)**
- django/db/models/sql/where.py
 - django/db/models/sql/datastructures.py
 - django/utils/regex_helper.py
 - django/utils/functional.py
 - django/db/models/fields/__init__.py

PR Quality Scorecard: django/django/pr19419

PR #19419 Quality Scorecard - django_django






Relevant Files

Relevant Files (5 total)

- django/utils/formats.py
- django/utils/translation/__init__.py
- django/utils/safestring.py
- django/utils/timezone.py
- django/template/__init__.py

PR Quality Scorecard: django/django/pr19445

PR #19445 Quality Scorecard - django_django



- Relevant Files
- Relevant Files (5 total)**
- django/db/models/fields/composite.py
 - django/db/models/expressions.py
 - django/db/models/utils.py
 - tests/model_regress/test_state.py
 - django/contrib/contenttypes/fields.py

Methodology

The data curation pipeline implements a multi-stage filtering approach inspired by the SWE-RL paper, focusing on extracting high-quality software engineering data from GitHub repositories. The pipeline consists of the following key components:

+ 1. Data Acquisition

- GitHub API integration for PR events and metadata
- Repository cloning for file content access
- Linked issue resolution and context gathering

2. Multi-Stage Filtering

- Bot and Automation Detection: Identifies and filters out automated PRs
- Size and Complexity Filtering: Ensures PRs are neither trivial nor unwieldy
- Content Relevance Filtering: Focuses on meaningful software engineering content

* 3. Relevant Files Prediction

- Identifies semantically related files not modified in the PR
- Uses import analysis and directory structure heuristics
- Enhances context for understanding code changes

@ 4. Quality Metrics Generation

The filtering pipeline maintains high precision by using progressive refinement, ensuring that only PRs with genuine software engineering value are retained while capturing detailed metadata about filtering decisions.

- Comprehensive quality scoring across multiple dimensions
- Metadata extraction for filtering decisions
- Relevance scoring based on problem-solving indicators